

Siddharth Surange

AI Engineer · Delivery Lead

ssurange.dev@gmail.com · sid-surange.github.io

CONTACT

Email

ssurange.dev@gmail.com

GitHub

github.com/SID-SURANGE

LinkedIn

linkedin.com/in/siddharthsurange

Site

sid-surange.github.io

SKILLS

Backend

Python, FastAPI, Flask, Pydantic, Docker, Swagger/Postman

LLMs & agents

OpenAI, Gemini, Claude, Llama, Qwen; agents, RAG, prompt design; LangChain, LlamaIndex, smolagents; Ollama, LM Studio

ML & CV

PyTorch, Hugging Face Transformers, spaCy, Pandas, NumPy; YOLO, RetinaNet, StyleGAN; NLP

Cloud & DevOps

Google Cloud: Vertex AI, Cloud Run, Cloud Storage, BigQuery; Terraform; GitHub Actions

UI

Gradio, Streamlit

EDUCATION

B.E. Information Technology

Acropolis Institute of Technology & Research, 2015

AI/ML Specialization

Springboard, 2019

CERTIFICATIONS

- Google Cloud Professional Data Engineer
- Oracle Generative AI Professional

AWARDS

- Transformer of the Month, Affine
- Spot Award, DXC Technology

WRITING (MEDIUM)

- Briefcast: a personal AI intelligence agent
- AgentForge: an AI agent with web search and image generation
- Deciphering the power of Vision Language Models
- What's new with OpenAI's gpt-4o-mini
- ProGAN, StyleGAN, StyleGAN2: exploring NVIDIA's breakthroughs

SUMMARY

AI engineer and delivery lead with 10+ years in software, 6 of them in AI and deep learning, and LLM work since 2021. Leads teams of 2–5 engineers building agentic and ML systems on top of real client systems, from scoping and design through production rollout. Comfortable across computer vision, NLP, LLM agents and Google Cloud.

EXPERIENCE

AI Engineer · Cognida

May 2025 – Present

- Concrete logistics automation** – led a team of 4 automating a building-materials client's order-assignment system with an agent-first approach, still relying on core optimization libraries.
- Sales Copilot** – led a team of 5 building an MS Teams sales bot that recommends quotes for Salesforce deals; backed by an agent within Snowflake that fetches quotes using the ML engine.

Senior Software Engineer · Tech Mahindra

Dec 2023 – May 2025

- Led analysis and design of optimized prompt strategies for Google Gemini's code-generation capabilities.
- Ran training sessions on prompt-engineering best practices, followed by a 50% increase in project delivery speed and higher team productivity.

ML Application Engineer · Egen

Dec 2021 – Nov 2023

- Medical entity extraction** – led a team of 2 on a cloud clinical-document dashboard for a healthcare provider. Terraform for infrastructure, GitHub Actions for CI/CD, Vertex AI and Longformer extracting clinical entities at 86% accuracy; ran 5+ model trials (spaCy and LLMs in the Vertex AI model garden) for selection.
- Cloud call/IVR migration** – led about 4 engineers on a reliable, scalable call/IVR integration using Google Cloud and Twilio call and speech APIs; consulted with the client to analyse their VoIP infrastructure and proposed the Twilio-based design.

Solution Leader · Brane Enterprises

Mar 2021 – Nov 2021

- Implemented pose estimation with DensePose on video data and generated 3D models in Blender/Unity to analyse patients' exercise routines.

Data Science Associate · Affine

Dec 2019 – Feb 2021

- Led end-to-end computer-vision object detection (RetinaNet, YOLO) for three enterprise clients: image analysis and annotation, training, optimization and inference, reaching 88% accuracy.
- Built a StyleGAN proof of concept generating realistic game-character images, folding in user-testing feedback to improve avatar diversity, delivered in 4 weeks.

Professional Product Developer · DXC Technology

Aug 2015 – Jun 2019

- Led an on-site Agile implementation in Sweden for a home-insurance application: ran daily ceremonies and cross-functional communication.
- Delivered 30+ new features as enhancements to a health-insurance app.

PERSONAL PROJECTS

Briefcast github.com/SID-SURANGE/briefcast

Automated Google AI briefing agent: pulls four Google blog feeds, summarises and ranks new posts, and composes a daily briefing served as a web page with an archive and source-health view. Two-layer dedup (URL hash + title-embedding similarity), top 8 items with mandatory citations, per-call cost logging. Scope was cut deliberately (RAG, Telegram, extra sources), each decision recorded in an ADR. FastAPI, PostgreSQL + pgvector, OpenRouter.

ResumeParser github.com/SID-SURANGE/ResumeParser

HR resume analysis that runs fully offline on local LLMs (8-bit quantized Llama 3.1/3.2, IBM Granite): extracts structured data from PDFs, flags missing sections, generates tailored interview questions. FastAPI, Gradio, IBM Docling, LM Studio.

PageSense · AgentForge github.com/SID-SURANGE/AI-Sandbox

PageSense: Chrome extension with a FastAPI/Qdrant backend for semantic search over browsing history. AgentForge: multi-tool agent with web search and image generation on Hugging Face smolagents. Also in the repo: a GitHub-repo chat app (LlamaIndex, spaCy) and a multimodal RAG barista assistant.

SELECTED CLIENT WORK

Medical Entity Extraction · Healthcare provider

- Extracted medical entities for a clinical-notes viewer; built the ML flow with spaCy, Vertex AI and transformer models.
- Owned end-to-end deployment and CI/CD; provisioned GCP resources as code with Terraform through a GitHub Actions continuous-deployment pipeline.

Twilio Call/IVR Integration · Telecom client

- Migrated an existing telecom network setup to Twilio; led the Twilio call integration into a GCP-hosted cloud application.
- Added audio and video calling for multiparty conference sessions; used Twilio IVR and advanced speech models to handle different IVR scenarios.

Natural Gas Source Detection · Offshore energy-exploration client

- Located gas and oil sources on sea beds. Classified gas presence in water-column images with a fast.ai ResNet34 model using Focal Loss.
- Added a detection model for exact plume coordinates; YOLOv4 with a ResNet101 backbone gave the best results.

Defect Detection · Consumer-goods manufacturer

- Detected product defects with state-of-the-art object detection: EDA and augmentation, RetinaNet with a suitable backbone, result analysis with the right metrics, and an end-to-end inference pipeline.